WF

SUBMISSION TO THE PANEL ON SECURITY · LEGISLATIVE COUNCIL OF HONG KONG

Wang Fuk Court Fire Inquiry: the record is public, the answers are hard to find

A citizen-built, machine-readable companion to the Independent Committee's published record

To	Members of the Panel on Security, Legislative Council of Hong Kong
From	Wenyun Fu, Lucia Luo and Simon Wang (Wang Fuk Court Fire Report project)
Date	19 October 2026
Status	**Draft for review** — 21 October 2026 · **not yet submitted**; to be sent by Lucia Luo once Simon has approved
Subject	Accessibility of the inquiry record ahead of the Committee's final report, due by end of October 2026

1 · The original site: what it is, how it is built, where it falls short

Background. On **26 November 2025** a fire swept Wang Fuk Court in Tai Po — an eight-block Home Ownership Scheme estate of some 4,600 residents, built in 1983 and then undergoing renovation. It killed **168 people**, including firefighter Ho Wai-ho, displaced thousands, and was the city's worst fire since 1948. On **12 December 2025** the Chief Executive appointed an **Independent Committee**, chaired by Mr Justice David Lok, to establish the facts and make recommendations. Its **seven Terms of Reference** cover the cause of the fire, the renovation works and their supervision, the adequacy of existing laws and penalties, the emergency response, and related questions. Six rounds of public hearings ran from March to July 2026; the Committee's report is due to the Chief Executive by the end of October 2026.
Links: the Committee's site · mandate & Terms of Reference · notices and press releases · the document archive

The site we started from. The Committee publishes its evidentiary record at ic-wangfukcourtfire.gov.hk: a home/"About" page (mandate, membership, Terms of Reference), a **Documents** page holding all 243 files, a **Notices and Press Releases** page, and a separate Chinese mirror. The files are grouped by *working stage* — notices, then each round of hearings, witness statements, expert reports, closing submissions — because the design logic is the Committee's own schedule: what it did, in the order it did it. Everything is free to download; nothing is arranged by the questions the public actually asks.

Where it falls short — the starting point of our work:

- **Filed by schedule, not by question.** To answer one question you must open dozens of files filed under different weeks.
- **Keyword search, not answers.** Searching “collusion” returns about 30 hits — hearing transcripts in date order, each mentioning the word in passing; none of them tells you whether collusion occurred (Figure 2).
- **347 pages in 46 documents are scanned images** — no text, so no search can reach them at all.
- **No summaries, no citations, no map.** A reader cannot see what a document says without reading it, nor trace a claim to a page.
- **Two languages, two separate archives** — no cross-language search, so Chinese readers and English readers meet different silos of the same record.



Figure 1 · The starting point: the Committee’s own site, 243 PDFs filed in the order the work was done

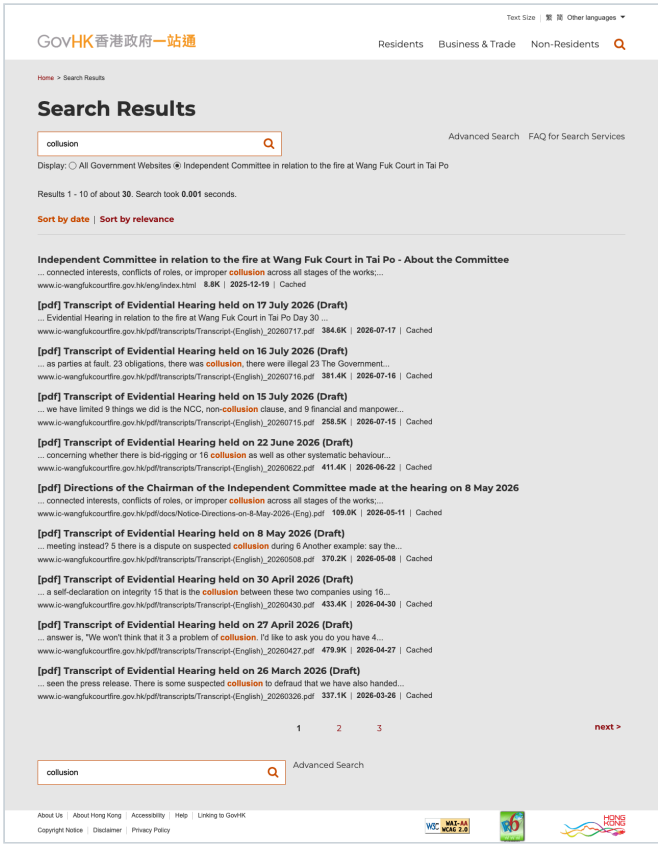


Figure 2 · The official search for “collusion”: about 30 hits — hearing transcripts in date order, none of them an answer

2 · What we built: a new site, vibe-coded with agentic AI

We did not write a new report — we rebuilt the archive as a **new website**, free and bilingual, at [wangfukcourtfirereport.simonsays.hk](). The whole pipeline — crawler, OCR, page converter, the tagging models, the site generators — was **written by an agentic AI working with us** ("vibe coding"): the agent wrote and debugged the code, we set the questions, checked the claims and verified every figure against the source PDF. That is why a volunteer team could do in **days** what a procurement cycle would take months to commission — and it is the same argument we make to the Panel in section 4: **if tools can do this, the publisher can do this at source.**

243 PDFs
crawled & audited · none missed

→

Plain text
6,377 pages · 347 OCR'd · 8,785 images

→

7 Terms of Reference
every paragraph tagged by meaning, not keyword

→

A new bilingual site
search · map · page citations · source PDFs

How the new site outsmarts the original — feature by feature:

What a reader needs	The original site	Our site
Organisation	filed by the Committee's work schedule	**grouped by the seven Terms of Reference** — the questions, not the calendar
Search	keyword over PDFs: "collusion" → ~30 raw transcript hits, no answer	**full text across 1.9m words**, meaning-tagged, EN + 中文 in one place
Scanned pages	347 pages in 46 documents are images — invisible to search	**OCR'd and searchable**
Each document	a raw PDF	**summary, key points, contents, page-level citations**, link to the source PDF
Connections	none — you must know what to open	**interactive document map** + one narrative page per Term of Reference
Languages	two separate archives	**one index, both languages** (中文站)
Machine-readable	no	**JSON indexes, static HTML, stable URLs** — usable by tools and by people
Cost to build	—	**volunteers + agentic AI, days**

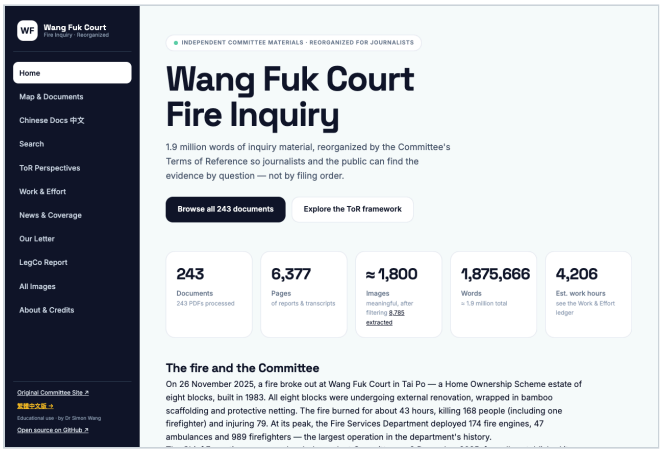


Figure 4 · Our home page: the same record re-organised around the seven questions the Committee was asked

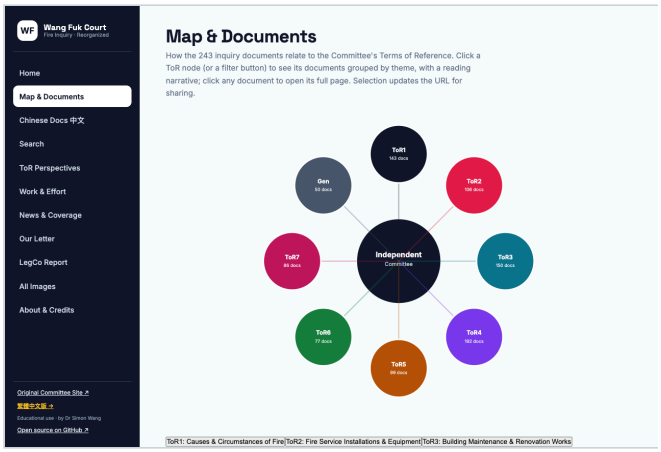


Figure 5 · Every document page: summary, key points, page-level citations and a link to the original PDF

What this gives a Member of the Panel — three worked examples, each with a live demo:

- **Example A — check a number before you ask it.** Suppose a colleague plans to ask about the 85 shutdown notices. Open search: shutdown → the 10 April transcript carries it 127 times, and the 8 April transcript has the contractor ticking *"Fire hydrant/fire hose reel system"* on the notice (Page 39). Thirty seconds, with a citable page.
- **Example B — test a policy idea against the record.** Thinking of proposing stronger debarment powers? search: suspend shows one contractor already *"ordered to suspend the registration for three*

weeks" (8 Apr) and Prestige *"fined and suspended from taking minor works by BD"* in 2017 (17 Jul) — i.e. existing powers exist but were thin.

- **Example C — brief staff or constituents bilingually.** One page per question in ToR narratives, mirrored in 繁體中文, each claim traced to a page and the original PDF — no 6,377-page archive required.
- **And when the report lands (due end-October):** our forecast of its findings — 11 predictions, each tied to the evidence — is published, with the evidence, on our forecast page; the ToR view is where each prediction will be marked confirmed or refuted, with page citations, within 24 hours of publication.
- **Freshness, demonstrated:** on 20 October we re-crawled the Committee site and found **56 transcripts revised on 15 October**; they were re-downloaded, re-processed, re-indexed and redeployed the same evening, and a re-audit returned *changed = 0*. The site tracks the record rather than lagging it (audit log on issue #4).

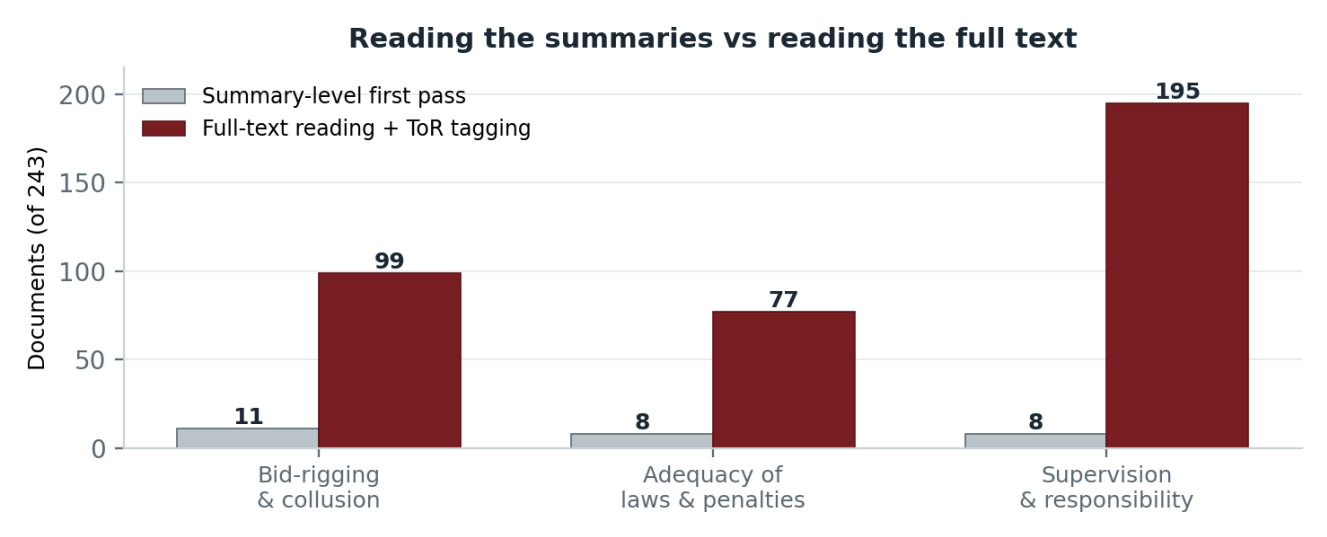


Figure 3 · Why full-text reading matters: documents found per Terms of Reference, summaries vs full text

3 · For journalists: source it, translate it, anticipate it

Reporters on deadline need three things the archive cannot give them: **a source they can link in one click**, **the other language**, and **prior art** — what was already known, so the story is not re-told from the press release. Our site supplies all three: every claim carries a page number and the source PDF; the same index exists in English and 繁體中文; and the machine-readable output lets a newsroom pull facts into its own database (we indexed 595 fire-day articles that way).

- **Case 1 — the tender story, sourced in one click.** The official search for "collusion" returns ~30 raw transcript hits and no answer; our Terms-of-Reference view returns **99 documents** — including a contractor that filed **85 shutdown-renewal notices without visiting the estate**, a renovation subcontract agreed **orally with no written contract**, and a contract attachment whose material-data folder was **empty**, none of which uses the words "bid-rigging". Demo: search: collusion · 30 Apr transcript (329 tender hits).

- **Case 2 — the Chinese story, explained.** Chinese-language coverage of 宏福苑 runs **1,004 articles against 106 English** in the same eight weeks, yet a Chinese outlet cannot cite the inquiry record: the Committee publishes Chinese PDFs, but filed by schedule and unsearchable, so a reporter must paraphrase. Our Chinese site fixes exactly that — take the 明報's 18 Oct story "怡樂投標現相同董事": the reporter can now attach the 30 April tender testimony and the ToR-5 list as evidence, in Chinese, with page numbers — a follow-up that no one else can write from the official site.
- **Case 3 — anticipating the final report.** The report is due by the end of October. We have published an 11-point forecast with the evidence behind each prediction — e.g. *"penalties are too low → legislation recommended"* rests on the 10 April testimony that the maximum penalty is *"a fine of $50,000, there is no imprisonment term"* while the Lifts and Escalators Ordinance is "much higher". A journalist can therefore pre-report the story, then publish the day the report lands with our ToR view open beside it — confirming or refuting each prediction with page-level citations. The same works for the **999 call** story of 15 October: search: 999 shows the issue was already examined in the 15 April transcript (105 hits) — prior art that turns "shocking new disclosure" into "the pattern the inquiry already saw".

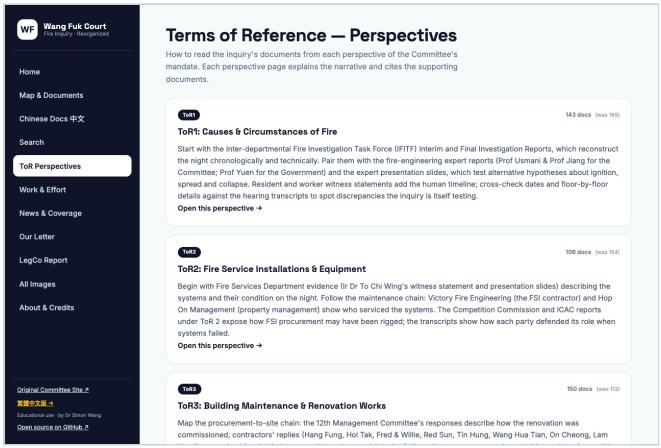


Figure 6 · ToR narratives: one page per question, in both languages



Figure 7 · The Chinese site: the same citations in 繁體中文

4 · The larger point: publish documents that can be used

Central argument: publication without usability is not accountability. When a public body releases a record, the test is whether a person — or a machine — can **question it on the day it appears**. Everything below follows from that one sentence.

- **1 · This is routine work.** Government publishes documents constantly — reports, papers, annexes, transcripts. The format is therefore a standing policy choice, not a one-off.
- **2 · The test is the format.** Is it machine-readable and agent-friendly — a structured file, **XML or CSV through data.gov.hk** — or a traditional, image-based format only a person can read?
- **3 · The Government already passes its own test.** The Chief Executive's Policy Address is deposited on data.gov.hk as **three XML files a year** (English, 繁體中文, 简体中文) with a data dictionary: 2025 · 2026 (checked 20 October 2026). The practice exists, it is routine, and it works.
- **4 · The inquiry record is not even that.** It is **243 documents, 6,377 pages, ~1.9 million words** — the largest body of evidence in this affair — released as **PDFs alone**, of which **347 pages**

in 46 documents are scanned images with no text at all; there is **no deposit on data.gov.hk**, no structured metadata, no page anchors, keyword-only search, and the English and Chinese files sit in two separate silos. A Policy Address gets three XML files; a public inquiry gets a pile of PDFs.

- **5 · The consequence is predictable.** A keyword box cannot find what no keyword contains, so the public gets the press release; and an AI tool pointed at such an archive reproduces the same failure, only more fluently.
- **6 · Our position.** Deposit as structured data — **XML or CSV through data.gov.hk — at the point of publication**: real text not page images, headings and metadata not a flat filing order, stable links and page anchors, both languages. And go further: arrange it around the **reader's questions**, so that people and machines can use it the first time. Volunteers with agentic AI rebuilt 6,377 pages in days — the barrier is habit, not capability, and it should not fall to volunteers to do the publisher's work afterwards.

A note on how this was made: this work is AI-assisted end to end — an agent wrote the crawler, the OCR and tagging pipeline and the site itself ("vibe coding"), while we set the questions and checked the claims. That is precisely the point: if tools can do this in days, they can do it at source. Every claim above links to the original document, so readers should check us against it.

Links

Our live site (EN + 繁體中文)	All 243 documents
Full-text search	Seven Terms of Reference
中文站 Chinese site	Our letter in the SCMP
News coverage analysis	Official Committee site
This report, as a web page	This report, as a PDF
Previous version (20 Oct archive)	Source code & data

The full record — summaries, search, seven ToR narratives — is live at:

https://wangfukcourtfirereport.simonsays.hk/engagement/legcoReport.html

Search: https://wangfukcourtfirereport.simonsays.hk/ · News: https://wangfukcourtfirereport.simonsays.hk/news.html

Kept current; the QR code always opens the latest version.

Scan to open the live report page

7 · Signatories

▼ **Change log** — versions of this report (times in HKT, UTC+8)

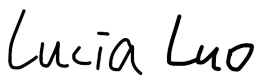
- **19 Oct 2026, 19:57** — first final draft rendered as PDF (4 pages) for submission to the Panel.
- **19 Oct 2026, 20:05** — dedicated page created at `/engagement/legcoReport.html`; clickable links and QR added; status set to final.
- **19 Oct 2026, 20:24** — Chinese-language coverage findings added (宏福苑: 1,004 articles vs 106 English).

- **19 Oct 2026, 20:51** — journalist-facing section added; draft email to the Panel Secretariat published at `/engagement/legcoEmail.html`.
- **19 Oct 2026, 20:57 – 21:05** — three signatures embedded (transparent PNGs), student roles added, HKBU disclaimer added.
- **19 Oct 2026, 21:15** — site navigation renamed "Engagement" → "LegCo Report"; the old `legcoDraft` page removed.
- **19 Oct 2026, 21:32** — restructured into four sections with figures, a chart and a flow diagram; text cut to 55%.
- **20 Oct 2026, 20:18** — previous version archived at `/engagement/legcoReport_20_Oct.html`.
- **21 Oct 2026, 19:10** — status corrected to **"Draft for review — not yet submitted"** (the report is still under review; Lucia will send only when Simon approves).
- **20 Oct 2026, 23:36** — forecast moved off GitHub onto a dated frontend page (`/engagement/forecast.html`); report links updated; page widened (side margins reduced).
- **20 Oct 2026, 22:35 — this revision**: §2 rewritten around the agentic-AI ("vibe coded") build with a feature-by-feature comparison table and three Member examples with live demos; §3 gains a clearer journalist value statement, a rewritten Case 2 and a new Case 3 (anticipating the final report); §4 restructured as a six-step chain with an explicit central argument and a detailed "not even that" bullet; previous version archived at `legcoReport_20_Oct_v2.html`.
- **20 Oct 2026, 21:25** — verified on data.gov.hk: the 2025 *and* 2026 Policy Addresses are each deposited as three XML files with a data dictionary; the hedge replaced with the fact, and the contrast sharpened.
- **20 Oct 2026, 21:20** — "machine-readable" defined precisely as a structured **XML/CSV deposit on data.gov.hk**; Policy Address example re-framed accordingly.
- **20 Oct 2026, 20:20** — revision: background overview added (the fire, the Committee, its Terms of Reference, with links); original site's structure and design logic described; critique kept as bullets; "new site" framing, figure-by-figure rationale, and a value list for Panel members; two journalist case studies with links; Policy Address example and the machine-readable position; AI-assisted note; this change log added; Figure 2 caption corrected (the official search returns ~30 hits, not "no results").



Wenyun Fu

Final-year student, Department of Journalism, Hong Kong Baptist University; co-author, "Fire inquiry documents should be more accessible", South China Morning Post, 19 August 2026



Lucia Luo

Third-year student, Department of Management, Marketing and Information Systems, Hong Kong Baptist University; contributor, Wang Fuk Court Fire Report project



Simon Wang, PhD

Lecturer in English, Hong Kong Baptist University; co-author, SCMP letter, 19 August 2026

Disclaimer: The views expressed in this paper are those of the authors. They do not represent the Hong Kong Baptist University and have not been endorsed by the University.